

Some content has been intentionally redacted (proprietary information and screenshots).

This documentation was created using docs-as-code, with a Static Site Generator (SSG) seamlessly converting Markdown content into HTML. The SSG took charge of formatting, while shortcodes were used for callouts and code blocks. Links between internal applications and documentation have been removed.

This snippet is a glimpse of an internal product's landing page; the full document includes Use Cases (with Mermaid diagrams), Tutorials and Guides, Security Guidance, and Support.

## Apache Cassandra Overview

Apache Cassandra is an open-source, distributed NoSQL database that uses a partitioned wide-column data model. It is designed for high availability and scalability, providing eventual consistency by default and configurable consistency guarantees through its tunable consistency model.

Cassandra enables applications to balance consistency, availability, and latency by selecting appropriate consistency levels per operation. Clusters can be deployed within a single region or across multiple geographic regions, depending on availability and resiliency requirements.



### Use Cases

#### **Recommendation Engines**

Serves as a high-scale datastore for recommendation systems that analyze user behavior and preferences.

#### **Real-Time Big Data Analytics**

Enables handling of large datasets and high-velocity write workloads, making it suitable for real-time analytics scenarios.

#### **Time-Series Data**

Allows for storing and querying time-ordered data such as metrics, logs, and sensor data.

#### **Location-Based Services**

Supports geospatial and location-based workloads, typically in combination with application-level querying and indexing strategies.

## How It Works

Cassandra's architecture is built for horizontal scalability and fault tolerance, allowing it to handle large volumes of concurrent read and write operations across distributed nodes.

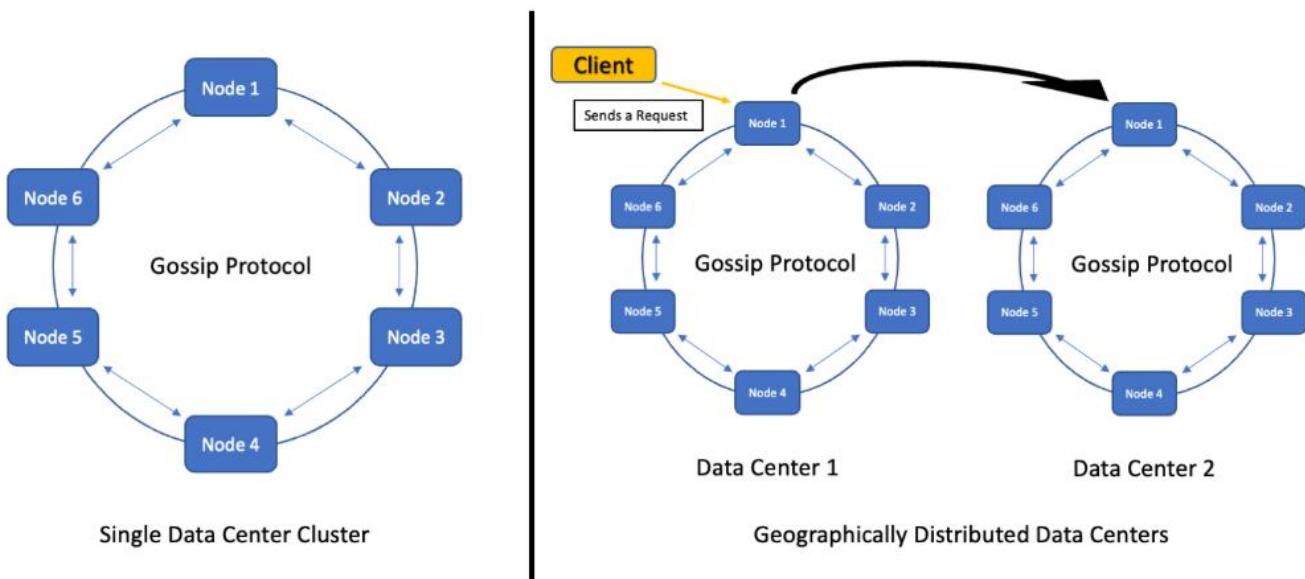
In multi-region or multi-data-center deployments, network latency can affect read and write performance. To address this, Cassandra provides configurable consistency levels—such as QUORUM—that allow applications to trade off latency and consistency based on their requirements and replication strategy.

For example, a replication factor of three ensures that each piece of data is stored on three different nodes, improving resilience and availability in the event of node or availability-zone failures.

The diagrams illustrate:

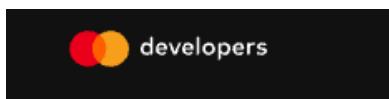
- A **single data-center Cassandra cluster**, where nodes communicate using the gossip protocol to share state information.
- A **geographically distributed deployment**, where multiple data centers operate together as part of a single logical cluster.

These diagrams are intended to demonstrate Cassandra's peer-to-peer architecture and replication model at a conceptual level.



## Tutorials and Guides

Here is a snapshot of the guides created for this product. Full details have been omitted, as they contain proprietary screenshots and code snippets.



## Apache Cassandra

### DOCUMENTATION

[Overview](#)

[Use Cases](#) ▾

[Tutorials and Guides](#) ^

[Onboarding Guide](#) ▾

[Naming Standards](#) ▾

[Cassandra Security](#) ▾

[Monitoring Guide](#) ▾

[Splunk Onboarding](#) ▾

[Support](#) ▾