



Some content has been intentionally redacted (proprietary information and screenshots).

This documentation was created using docs-as-code, with a Static Site Generator (SSG) seamlessly converting Markdown content into HTML. The SSG took charge of formatting, while shortcodes were used for callouts and code blocks. Links between internal applications and documentation have been removed.

This snippet is a glimpse of an internal product's landing page; the full document includes Use Cases (with Mermaid diagrams), Tutorials and Guides, and Support.

Product Onboarding Assistant

The Product Onboarding Assistant (POA) is a generative AI-powered solution designed to accelerate product onboarding and integration. It supports customer-facing teams by providing fast, accurate answers to common setup and configuration questions across multiple stages, including planning, implementation, warranty, and troubleshooting.

POA helps explain products and implementation details in clear, accessible language. By simplifying technical concepts and API guidance, it improves the clarity and consistency of information, helping reduce onboarding time, accelerate time to value, and support a smoother customer experience.



Key Capabilities

Conversational Assistant

POA provides a user-centric chat experience that enables users to quickly find relevant information.

Key capabilities include:

- Answers grounded in the latest documentation using Retrieval-Augmented Generation (RAG)
- Source citations that link responses back to reference materials
- Feedback mechanisms to rate answers and improve quality
- The ability for authorized users to refine responses for clarity
- An embedded, in-page chatbot experience
- Integration with enterprise productivity and analytics tools

Administrative Controls

POA includes administrative capabilities that support quality, governance, and continuous improvement.

Key capabilities include:

- Human-in-the-loop review workflows to evaluate and refine responses
- Visibility into response feedback and historical revisions
- Management of documentation and knowledge sources to keep information current
- Reporting and monitoring dashboards to track usage and quality trends


How It Works

POA uses a large language model (LLM combined with Retrieval-Augmented Generation (RAG) to produce responses grounded in authoritative documentation and curated knowledge sources.

Relevant content is retrieved at query time and provided to the model as context, helping ensure responses are accurate, up to date, and traceable to source material.

Human reviewers can refine generated responses for accuracy and clarity, with additional review steps applied as part of a controlled quality-improvement process. Feedback is used to improve content, retrieval logic, prompts, and model configurations through governed update cycles. Over time, this approach improves response quality while maintaining consistency and oversight.

The following diagram illustrates the system architecture at a high level:



Architecture Overview

Roles

Role	Description
Domain User	Users who interact with the assistant by asking questions and receiving answers.
Domain SME User	Subject Matter Experts (SMEs) who review and refine content to ensure accuracy.
Domain Admin User	Manages knowledge sources, review workflows, and overall configuration.

Data and Processing Flow

Component	Description
Customer Q&A	User questions and system-generated responses.
Peer Review and Refine	Expert review processes that improve response quality.
Customer Q&A (Refined)	Reviewed and refined content that meets quality standards.
Product Docs and Metadata	Source materials used to provide contextual knowledge.
Upload, Chunk, and Embed	Processes that prepare content for semantic retrieval.
Vectors	Embedded representations of content used for similarity search.
Vector Index	A searchable index that enables efficient semantic retrieval.
Chat Q&A	The interaction layer that delivers responses to users.

Models and AI Components

Component	Description
Embedding Model	Generates vector embeddings from content.
Summarization Model	Distills and summarizes key information from documentation.
RAG Pipeline	Retrieves relevant context and augments queries prior to response generation.
Foundation Model	Underlying language model supporting generation and reasoning.
Fine-Tuning	Controlled processes for updating models using curated feedback and data.

Tutorials and Guides

Here is a snapshot of the guides created for this product. Full details have been omitted, as they contain proprietary screenshots and code snippets.

